

# Project topic:TensorFlow and Neural Networks Applications in HealthCare

**Student Name: Liangliang Zhang**

## Problem Statement:

Use TensorFlow to build a Dense Neural Network that will be used to automatically classify fetal cardiotocogram to different fetal state (N, S, P) based on their diagnostic features data provided by the UCI Machine Learning Repository.

## Overview of Technology

TensorFlow and TensorBoard was used to built Multilayer Dense Neural Network model and monitor loss function on training dataset.

TensorFlow is an open source software library for numerical computation using data flow graphs. Nodes in the graph represent mathematical operations, while the graph edges represent the multidimensional data arrays (tensors) communicated between them. The flexible architecture allows you to deploy computation to one or more CPUs or GPUs in a desktop, server, or mobile device with a single API. (<https://www.tensorflow.org/> (<https://www.tensorflow.org/>))

## Descreption of Data

2126 fetal cardiotocograms (CTGs) were automatically processed and the respective diagnostic features measured. The CTGs were also classified by three expert obstetricians and a consensus classification label assigned to each of them. Classification was both with respect to a morphologic pattern (A, B, C. ...) and to a fetal state (N=normal; S=suspect; P=pathologic).

URL: <http://archive.ics.uci.edu/ml/machine-learning-databases/00193/>  
(<http://archive.ics.uci.edu/ml/machine-learning-databases/00193/>)

Size: 1.66 MB., sample size: 2130

Format of data file: .xls file of Microsoft Excel

## Hardware

Windows PC with Intel Core M-5Y10c CPU (0.8GHz, 998MHz) and 4GB RAM

## Sofeware

Anaconda with Python 3.6.1

TensorFlow 1.3.0 <https://pypi.python.org/pypi/tensorflow/1.3.0>  
(<https://pypi.python.org/pypi/tensorflow/1.3.0>)

## Lessons learned & Pros/Cons

After tuning, my final Neural Network model gives a prediction accuracy of ~92% in training data and a prediction accuracy of ~90% in validation data. This model performs reasonably well and I suppose that if we have more observations, especially observations of the minority class, we could have built a more powerful neural network.

## YouTube URLs:

short (2 min): <https://www.youtube.com/watch?v=dwUjhR7LHFY> (<https://www.youtube.com/watch?v=dwUjhR7LHFY>)

long (15 min): [https://www.youtube.com/watch?v=25v\\_I7LKyBU](https://www.youtube.com/watch?v=25v_I7LKyBU) ([https://www.youtube.com/watch?v=25v\\_I7LKyBU](https://www.youtube.com/watch?v=25v_I7LKyBU))

```
In [1]: # import libraries
import tensorflow as tf
import numpy as np
import pandas as pd
```

## Steps & Demonstration

### 1. load data

Data fiel CTG.xls was downloaded from <http://archive.ics.uci.edu/ml/machine-learning-databases/00193/> (<http://archive.ics.uci.edu/ml/machine-learning-databases/00193/>)

```
In [2]: #Load data
ctg = pd.read_excel('CTG.xls', sheetname = 2)
```

### 2. Data cleaning

```
In [3]: #check data shape
print ('CTG data shape:', ctg.shape)
#check data head
ctg.head()
```

CTG data shape:, (2130, 40)

Out[3]:

	FileName	Date	SegFile	b	e	LBE	LB	AC	FM	UC	...	C	D
0	NaN	NaT	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN
1	Variab10.txt	1996-12-01	CTG0001.txt	240.0	357.0	120.0	120.0	0.0	0.0	0.0	...	0.0	0.0
2	Fmcs_1.txt	1996-05-03	CTG0002.txt	5.0	632.0	132.0	132.0	4.0	0.0	4.0	...	0.0	0.0
3	Fmcs_1.txt	1996-05-03	CTG0003.txt	177.0	779.0	133.0	133.0	2.0	0.0	5.0	...	0.0	0.0
4	Fmcs_1.txt	1996-05-03	CTG0004.txt	411.0	1192.0	134.0	134.0	2.0	0.0	6.0	...	0.0	0.0

5 rows × 40 columns



```
In [4]: #check data tail
ctg.tail()
```

Out[4]:

	FileName	Date	SegFile	b	e	LBE	LB	AC	FM	UC	...	C
2125	S8001045.dsp	1998-06-06	CTG2127.txt	1576.0	3049.0	140.0	140.0	1.0	0.0	9.0	...	0.0
2126	S8001045.dsp	1998-06-06	CTG2128.txt	2796.0	3415.0	142.0	142.0	1.0	1.0	5.0	...	0.0
2127	NaN	NaT	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN
2128	NaN	NaT	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN
2129	NaN	NaT	NaN	NaN	NaN	NaN	NaN	NaN	564.0	23.0	...	NaN

5 rows × 40 columns



Drop the column of Filename, Date and SegFile, as these information has absolutely no predictive power for determining the state of a CTG image. Keeping these information will only confuse our model when training neural network.

```
In [5]: ctg.drop(['FileName', 'Date', 'SegFile'], axis = 1, inplace = True)
```

Drop first row as it is blank, then drop a few rows from the bottom as they contain meaningless information.

```
In [6]: ctg_clear = ctg.drop(ctg.index[[0, 2127, 2128, 2129]])
```

Check if there are any missing values.

```
In [7]: print ('Having missing values? :', ctg_clear.isnull().any().any())
```

Having missing values? : False

Let's check the shape and statistical descriptions after cleaning:

```
In [8]: print ('Data shape after cleaning', ctg_clear.shape)
```

Data shape after cleaning (2126, 37)

```
In [33]: ctg_clear.head()
```

Out[33]:

	b	e	LBE	LB	AC	FM	UC	ASTV	MSTV	ALTV	...	C	D	E	AD	DE	LD
1	240.0	357.0	120.0	120.0	0.0	0.0	0.0	73.0	0.5	43.0	...	0.0	0.0	0.0	0.0	0.0	0.0
2	5.0	632.0	132.0	132.0	4.0	0.0	4.0	17.0	2.1	0.0	...	0.0	0.0	0.0	1.0	0.0	0.0
3	177.0	779.0	133.0	133.0	2.0	0.0	5.0	16.0	2.1	0.0	...	0.0	0.0	0.0	1.0	0.0	0.0
4	411.0	1192.0	134.0	134.0	2.0	0.0	6.0	16.0	2.4	0.0	...	0.0	0.0	0.0	1.0	0.0	0.0
5	533.0	1147.0	132.0	132.0	4.0	0.0	5.0	16.0	2.4	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0

5 rows × 37 columns




```
In [10]: ctg_clear.describe()
```

Out[10]:

	b	e	LBE	LB	AC	FM	UC
count	2126.000000	2126.000000	2126.000000	2126.000000	2126.000000	2126.000000	2126.000000
mean	878.439793	1702.877234	133.303857	133.303857	2.722484	7.241298	3.659925
std	894.084748	930.919143	9.840844	9.840844	3.560850	37.125309	2.847094
min	0.000000	287.000000	106.000000	106.000000	0.000000	0.000000	0.000000
25%	55.000000	1009.000000	126.000000	126.000000	0.000000	0.000000	1.000000
50%	538.000000	1241.000000	133.000000	133.000000	1.000000	0.000000	3.000000
75%	1521.000000	2434.750000	140.000000	140.000000	4.000000	2.000000	5.000000
max	3296.000000	3599.000000	160.000000	160.000000	26.000000	564.000000	23.000000

8 rows × 37 columns



### 3. Extract feature and lables

The dataset has two types of labels: morphologic pattern and fetal state. In this project, we only use fetal state label to perform a 3-class classification. Then fetal labels was then onehot encoded to dummy variables.

```
In [34]: features = ctg_clear.iloc[:, :-12].values
labels = ctg_clear.iloc[:, -1].values

labels_onehot = pd.get_dummies(labels)

print ('Number of observations:', features.shape[0])
print ('Number of features:', features.shape[1])
print ('number of labels:', labels_onehot.shape[1])
```

```
Number of observations: 2126
Number of features: 25
number of labels: 3
```

#### 4. Train and validation data split

Then entire dataset was randomly split into training (80% 1700 cases) and validation (20% 426 cases) dataset. Train dataset is used for training our neural network, and validation dataset is used for testing the accuracy of our model.

```
In [184]: from sklearn.model_selection import train_test_split
# Take 1/5 images from the training data, and leave the remainder in training
train_dataset, valid_dataset, train_labels, valid_labels = train_test_split(features, labels,
train_labels, valid_labels)
print('Training data/label shape: ', train_dataset.shape, train_labels.shape)
print('Validation data/label shape: ', valid_dataset.shape, valid_labels.shape)
```

```
Training data/label shape: (1700, 25) (1700, 3)
Validation data/label shape: (426, 25) (426, 3)
```

```
In [185]: #check the propotion of each class in train and validation data
print ('Propotion for each class in train data:', np.sum(train_labels, axis=0)/train_labels.shape[0])
print ('Propotion for each class in validaion data:', np.sum(valid_labels, axis=0)/valid_labels.shape[0])
```

```
Propotion for each class in train data: [ 0.77411765  0.14117647  0.08470588]
Propotion for each class in validaion data: [ 0.79577465  0.12910798  0.07511737]
```

The propotion for each class is similar in training and validation dataset, so we will have all information needed in training data.

#### 5. Dense Neural Network (DNN) model

##### 5.1 Define a few useful functions

```
In [186]: # calculate accuracy by identifying validation cases where the model's highest-probability prediction matches the labels
def accuracy(predictions, labels):
    correct_prediction = tf.equal(tf.argmax(predictions, 1), tf.argmax(labels, 1))
    accuracy_pct = tf.reduce_mean(tf.cast(correct_prediction, tf.float32)) * 100.0
    #another way to calculate this is to use np like following
    return (100.0 * np.sum(np.argmax(predictions, 1) == np.argmax(labels, 1)) / predictions.shape[0])
    #return accuracy_pct.eval()
```

```
In [187]: def weight_variable(shape, name):
    initial = tf.truncated_normal(shape, stddev=1e-4)
    #initial = tf.truncated_normal(shape, stddev=np.sqrt(2.0/shape[0]))
    return tf.Variable(initial, name=name)

def bias_variable(shape, name):
    #initial = tf.constant(0.1, shape=shape)
    initial = tf.zeros(shape)
    return tf.Variable(initial, name=name)

split_by_half = lambda x,k : int(x/2**k)
```

## 5.2 Simple 2-layer DNN model with GradientDescentOptimizer

```

In [188]: valid_dataset = valid_dataset.astype(np.float32)
n_labels = 3
batch_size = 99
flattened_size = train_dataset.shape[1]
hidden_nodes = 100

graph = tf.Graph()
with graph.as_default():

    # Input data.
    tf_train_dataset = tf.placeholder(tf.float32, shape=(batch_size, flattened_size))
    tf_train_labelset = tf.placeholder(tf.float32, shape=(batch_size, n_labels))
    tf_valid_dataset = tf.constant(valid_dataset, name="ValidationData")

    # Variables.
    layer1_weights = tf.Variable(tf.truncated_normal([flattened_size, hidden_nodes]))
    layer1_biases = tf.Variable(tf.zeros([hidden_nodes]), name="biases1")
    layer2_weights = tf.Variable(tf.truncated_normal([hidden_nodes, n_labels]))
    layer2_biases = tf.Variable(tf.ones([n_labels]), name="biases2")

    # Model.
    def model(data, name):
        with tf.name_scope(name) as scope:
            layer1 = tf.add(tf.matmul(data, layer1_weights), layer1_biases, name="layer1")
            hidden1 = tf.nn.relu(layer1, name="relu1")
            layer2 = tf.add(tf.matmul(hidden1, layer2_weights), layer2_biases, name="layer2")
            return layer2

    # Training computation.
    logits = model(tf_train_dataset, name="logits")
    # Loss function
    loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=logits, labels=tf_train_labelset))

    # Optimizer.
    optimizer = tf.train.GradientDescentOptimizer(0.05).minimize(loss)

    # Predictions for the training, validation
    train_prediction = tf.nn.softmax(logits)
    valid_prediction = tf.nn.softmax(model(tf_valid_dataset, name="validation"))

```

```
In [189]: # define run model function
def run_session(num_epochs, name):
    with tf.Session(graph=graph) as session:
        tf.global_variables_initializer().run()
        merged = tf.summary.merge_all()
        writer = tf.summary.FileWriter("tmp/tensorflowlogs", session.graph)
        print("Initialized model:", name)
        for epoch in range(num_epochs):
            offset = (epoch * batch_size) % (train_labels.shape[0] - batch_size)
            batch_data = train_dataset[offset:(offset + batch_size), :]
            batch_labels = train_labels[offset:(offset + batch_size), :]
            feed_dict = {tf_train_dataset : batch_data, tf_train_labelset : batch_labels}
            _, l, predictions = session.run([optimizer, loss, train_prediction], feed_dict)
            if (epoch % 500 == 0):
                print('Minibatch loss at epoch %d: %f' % (epoch, l))
                print('Minibatch accuracy: %.1f%%' % accuracy(predictions, batch_labels))
                print('Validation accuracy: %.1f%%' % accuracy(valid_prediction, valid_labels))
```

```
In [179]: run_session(5001, "DNN_2layer")
```

```
Initialized model: DNN_2layer
Minibatch loss at epoch 0: 10123.847656
Minibatch accuracy: 14.1%
Validation accuracy: 75.6%
Minibatch loss at epoch 500: 0.536054
Minibatch accuracy: 84.8%
Validation accuracy: 75.6%
Minibatch loss at epoch 1000: 0.751381
Minibatch accuracy: 73.7%
Validation accuracy: 75.6%
Minibatch loss at epoch 1500: 0.622334
Minibatch accuracy: 79.8%
Validation accuracy: 75.6%
Minibatch loss at epoch 2000: 0.629530
Minibatch accuracy: 80.8%
Validation accuracy: 75.6%
Minibatch loss at epoch 2500: 0.752382
Minibatch accuracy: 74.7%
Validation accuracy: 75.6%
Minibatch loss at epoch 3000: 0.658178
Minibatch accuracy: 78.8%
Validation accuracy: 75.6%
Minibatch loss at epoch 3500: 0.616110
Minibatch accuracy: 79.8%
Validation accuracy: 75.6%
Minibatch loss at epoch 4000: 0.674650
Minibatch accuracy: 76.8%
Validation accuracy: 75.6%
Minibatch loss at epoch 4500: 0.668530
Minibatch accuracy: 76.8%
Validation accuracy: 75.6%
Minibatch loss at epoch 5000: 0.686899
Minibatch accuracy: 76.8%
Validation accuracy: 75.6%
```

After 5000 epoches, both training and validation accuracies are around 76%, which is similar to a



blind guess of first class.

Next, we modify several parameters of our DNN model to see if we can improve model performance. Modifications are listed below:

1. More hidden layers
2. Regularization and dropout to avoid over fitting
3. Alternative optimizer

Also, summary for loss function, train accuracy and validation accuracy were added to TensorBoard, so we can keep tracking our model performance.

### **5.3 4-layer DNN model with regularization, dropout and AdamOptimizer**

```

In [194]: batch_size = 340
          flattened_size = train_dataset.shape[1]
          hidden_nodes = 512
          lamb_reg = 0.001
          learning_rate = 0.001 # Learning rate for the momentum optimizer

          graph = tf.Graph()
          with graph.as_default():

              # Input data.
              tf_train_dataset = tf.placeholder(tf.float32, shape=(batch_size, flattened_size))
              tf_train_labelset = tf.placeholder(tf.float32, shape=(batch_size, n_labels))
              tf_valid_dataset = tf.constant(valid_dataset, name="ValidationData")
              tf_valid_labelset = tf.constant(valid_labels, name="ValidationLabels")

              # Variables.
              layer1_weights = weight_variable([flattened_size, hidden_nodes], name="weight1")
              layer1_biases = bias_variable([hidden_nodes], name="biases1")
              layer2_weights = weight_variable([hidden_nodes, split_by_half(hidden_nodes,1)], name="weight2")
              layer2_biases = bias_variable([split_by_half(hidden_nodes,1)], name="biases2")
              layer3_weights = weight_variable([split_by_half(hidden_nodes,1), split_by_half(hidden_nodes,2)], name="weight3")
              layer3_biases = bias_variable([split_by_half(hidden_nodes,2)], name="biases3")
              layer4_weights = weight_variable([split_by_half(hidden_nodes,2), n_labels], name="weight4")
              layer4_biases = bias_variable([n_labels], name="biases4")

              keep_prob = tf.placeholder("float", name="keep_prob")

              def model(data, name, proba=keep_prob):
                  with tf.name_scope(name) as scope:
                      layer1 = tf.add(tf.matmul(data, layer1_weights), layer1_biases, name="layer1")
                      hidden1 = tf.nn.dropout(tf.nn.relu(layer1), proba, name="dropout1")
                      layer2 = tf.add(tf.matmul(hidden1, layer2_weights), layer2_biases, name="layer2")
                      hidden2 = tf.nn.dropout(tf.nn.relu(layer2), proba, name="dropout2")
                      layer3 = tf.add(tf.matmul(hidden2, layer3_weights), layer3_biases, name="layer3")
                      hidden3 = tf.nn.dropout(tf.nn.relu(layer3), proba, name="dropout3")
                      layer4 = tf.add(tf.matmul(hidden3, layer4_weights), layer4_biases, name="layer4")
                      return layer4

              # Training computation.
              logits = model(tf_train_dataset, "logits", keep_prob)
              loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=logits, labels=tf_train_labelset))
              regularizers = (tf.nn.l2_loss(layer1_weights) + tf.nn.l2_loss(layer1_biases) +
                              tf.nn.l2_loss(layer2_weights) + tf.nn.l2_loss(layer2_biases) +
                              tf.nn.l2_loss(layer3_weights) + tf.nn.l2_loss(layer3_biases) +
                              tf.nn.l2_loss(layer4_weights) + tf.nn.l2_loss(layer4_biases))

              # Add the regularization term to the loss.
              loss += lamb_reg * regularizers
              #loss = tf.reduce_mean(loss + lamb_reg * regularizers)

              # Optimizer
              #global_step = tf.Variable(0, name="globalstep") # count number of steps taken
              #optimizer = tf.train.MomentumOptimizer(learning_rate=learning_rate, momentum=0.9).minimize(loss)
              optimizer = tf.train.AdamOptimizer(learning_rate=0.001, epsilon=1e-04).minimize(loss)

              # Predictions for the training, validation, and test data.
              train_prediction = tf.nn.softmax(logits)

```

```

valid_prediction = tf.nn.softmax(model(tf_valid_dataset, "validation", 1.0))
#saver = tf.train.Saver() # a saver variable to save the model

# acuuracy for training data
train_correct_prediction = tf.equal(tf.cast(tf.argmax(logits, 1), tf.float32), tf.cast(tf.argmax(valid_prediction, 1), tf.float32))
accuracy_train = tf.reduce_mean(tf.cast(train_correct_prediction, tf.float32))
# acuuracy for validation data
valid_correct_prediction = tf.equal(tf.cast(tf.argmax(model(tf_valid_dataset, "validation", 1.0), 1), tf.float32), tf.cast(tf.argmax(valid_prediction, 1), tf.float32))
accuracy_valid = tf.reduce_mean(tf.cast(valid_correct_prediction, tf.float32))

```

```

In [195]: def run_session_2(num_epochs, name, k_prob=1.0):

    with tf.Session(graph=graph) as session:
        tf.global_variables_initializer().run()

        # summaries
        loss_summary = tf.summary.scalar('Loss', loss)

        train_accuracy_summary = tf.summary.scalar('train_accuracy', accuracy_train)
        valid_accuracy_summary = tf.summary.scalar('valid_accuracy', accuracy_valid)

        merged = tf.summary.merge_all()
        writer = tf.summary.FileWriter("tmp/tensorflowlogs_3", session.graph)

        print('Initialized model:', name, "\n")
        for epoch in range(num_epochs):
            offset = (epoch * batch_size) % (train_labels.shape[0] - batch_size)
            batch_data = train_dataset[offset:(offset + batch_size), :]
            batch_labels = train_labels[offset:(offset + batch_size), :]
            feed_dict = {tf_train_dataset : batch_data, tf_train_labelset : batch_labels}
            _, l, predictions = session.run([optimizer, loss, train_prediction], feed_dict=feed_dict)
            writer.add_summary(loss_summary.eval(feed_dict=feed_dict), epoch)
            writer.add_summary(train_accuracy_summary.eval(feed_dict=feed_dict), epoch)
            writer.add_summary(valid_accuracy_summary.eval(feed_dict=feed_dict), epoch)
            #writer.add_summary(learning_rate_summary.eval(), epoch)
            if (epoch % 500 == 0):
                print("Minibatch loss at epoch {}: {}".format(epoch, l))
                print("Minibatch accuracy: {:.1f}".format(accuracy(predictions, batch_labels)))
                print("Validation accuracy: {:.1f}\n".format(accuracy(valid_predictions, valid_labels)))
            #save_path = saver.save(session, "tmp/" + name + ".ckpt")
            #print("Model saved in file: %s" % save_path)

```

```
In [196]: run_session_2(5001, "DNN_4layer_Adam", 1.0)
```

```
Initialized model: DNN_4layer_Adam
```

```
Minibatch loss at epoch 0: 1.098612666130066
```

```
Minibatch accuracy: 80.6
```

```
Validation accuracy: 79.6
```

```
Minibatch loss at epoch 500: 0.22571192681789398
```

```
Minibatch accuracy: 92.4
```

```
Validation accuracy: 89.0
```

```
Minibatch loss at epoch 1000: 0.17441688477993011
```

```
Minibatch accuracy: 94.1
```

```
Validation accuracy: 87.1
```

```
Minibatch loss at epoch 1500: 0.13863994181156158
```

```
Minibatch accuracy: 95.6
```

```
Validation accuracy: 88.0
```

```
Minibatch loss at epoch 2000: 0.06984758377075195
```

```
Minibatch accuracy: 98.5
```

```
Validation accuracy: 88.3
```

```
Minibatch loss at epoch 2500: 0.1138841062784195
```

```
Minibatch accuracy: 97.1
```

```
Validation accuracy: 88.3
```

```
Minibatch loss at epoch 3000: 0.04885120317339897
```

```
Minibatch accuracy: 99.7
```

```
Validation accuracy: 88.7
```

```
Minibatch loss at epoch 3500: 0.04376523569226265
```

```
Minibatch accuracy: 99.4
```

```
Validation accuracy: 89.9
```

```
Minibatch loss at epoch 4000: 0.05714946240186691
```

```
Minibatch accuracy: 99.1
```

```
Validation accuracy: 89.9
```

```
Minibatch loss at epoch 4500: 0.0718587189912796
```

```
Minibatch accuracy: 98.2
```

```
Validation accuracy: 88.7
```

```
Minibatch loss at epoch 5000: 0.035477880388498306
```

```
Minibatch accuracy: 99.7
```

```
Validation accuracy: 90.8
```

```
In [145]: run_session_2(5001, "DNN_4layer_Adam", 1.0)
```

```
Initialized model: DNN_4layer_Adam
```

```
Minibatch loss at epoch 0: 1.098612666130066
```

```
Minibatch accuracy: 81.0
```

```
Validation accuracy: 76.8
```

```
Minibatch loss at epoch 500: 0.3719613552093506
```

```
Minibatch accuracy: 88.0
```

```
Validation accuracy: 84.5
```

```
Minibatch loss at epoch 1000: 0.2221064567565918
```

```
Minibatch accuracy: 89.0
```

```
Validation accuracy: 87.1
```

```
Minibatch loss at epoch 1500: 0.24829697608947754
```

```
Minibatch accuracy: 88.0
```

```
Validation accuracy: 87.6
```

```
Minibatch loss at epoch 2000: 0.28749698400497437
```

```
Minibatch accuracy: 88.0
```

```
Validation accuracy: 88.3
```

```
Minibatch loss at epoch 2500: 0.2003413736820221
```

```
Minibatch accuracy: 94.0
```

```
Validation accuracy: 86.9
```

```
Minibatch loss at epoch 3000: 0.10720705986022949
```

```
Minibatch accuracy: 97.0
```

```
Validation accuracy: 89.2
```

```
Minibatch loss at epoch 3500: 0.17392706871032715
```

```
Minibatch accuracy: 94.0
```

```
Validation accuracy: 90.4
```

```
Minibatch loss at epoch 4000: 0.19757121801376343
```

```
Minibatch accuracy: 91.0
```

```
Validation accuracy: 89.0
```

```
Minibatch loss at epoch 4500: 0.12272512167692184
```

```
Minibatch accuracy: 97.0
```

```
Validation accuracy: 90.6
```

```
Minibatch loss at epoch 5000: 0.08677855879068375
```


```
Minibatch accuracy: 99.0
```

```
Validation accuracy: 89.4
```

Modified Neural Network model gives a prediction accuracy of ~99% in training data and a prediction accuracy of ~90% in validation data. Visualization from TensroBoard is shown below:

```
In [1]: from IPython.display import Image  
Image("TensorBoard.png")
```

Out[1]:



Different optimizer (MomentumOptimizer, AdamOptimizer, GradientDescentOptimizer), learning rate (0.0001, 0.001, 0.01, 0.1) and keep probability (1.0, 0.8, 0.5) were tested. The final DNN model (AdamOptimizer, learning rate =0.001, keep probability=1.0), which has the best performance on validation data, was shown above.

## 6. Conclusion

Our final Neural Network model gives a prediction accuracy of ~92% in training data and a prediction accuracy of ~90% in validation data. This model performs reasonably well and I suppose that if we have more observations, especially observations of the minority class, we could have built a more powerful neural network.